MAC · MEETINGS · OFF BY DEFAULT

Live Radar and Follow-through

Live Radar catches commitments, decisions, open questions and conflicts with earlier meetings while you record. Follow-through then plans the work for one of your actions and drafts each step for your approval. Both read with the Think model you chose in Settings → Models.

DURING THE MEETING (MADE-UP EXAMPLE)

- **Commitment** · "I'll send the revised pricing deck by Friday."
- **Decision** · "Let's keep the Team plan at five seats."
- **Conflict** · Launch "the week after", but on 19 Oct you agreed Monday the 3rd.
- **Question** · "What was conversion in September?" Answered from your screen, with the line it rests on.

AFTER THE MEETING

- **Follow Through…** on one action
- **Proposes** · an email draft, a calendar invitation, a task for Claude Code or Codex, or marking the action done
- **You approve each step** · edit it first if you like
- **Opens as a draft** · in your mail app, Calendar or coding agent. LokalBot never sends anything itself.

TURN IT ON

What	Where
Live Radar	Settings → Meetings → Live Radar → **Offer Live Radar during meetings**. Then press **Start Radar** during a recording; it runs for that recording only. A small panel can float over your meeting app, and is left out of screen sharing.
Screen answers, live answers	Two more switches in the same section, both off. Screen answers read your screen memory; live answers also read the folders you add there.
Follow-through	No setting. Choose **Follow Through…** on an action.
Turn-by-turn live text	Settings → Meetings → Processing. The live preview transcribes one turn at a time, about half a second to a second behind the speaker, where it otherwise takes about two and up to five. Live Radar then reads after each turn.

WHAT IT READS, KEEPS AND SENDS

- With a Think model on this Mac, nothing here leaves it. With a Think server or plan, it goes only to the origin you approved for Think, and revoking that approval stops it.
- **Live Radar** gives the model the new preview lines, up to 24 earlier lines, the meeting title, invited names, and up to six passages local search found in earlier meetings. Audio, speaker separation and search stay on this Mac. Email addresses are never sent.
- Every item must quote a line that was said, and an answer must quote the source it cites. Anything else is dropped. Findings are saved with the meeting and deleted with it.
- **Screen answers** send text read from up to three crops of earlier captures, never the pictures. **Live answers** never read hidden files, `node_modules`, secret-looking files, or files git doesn't track.
- **Follow-through** gets the action, its source lines, the meeting's summary and what its read tools return from your meetings. It can't read screen memory, and email addresses are filled in on this Mac. Full details are in `PRIVACY.md`.

LIMITS

- Mac only. Live Radar reads about every 20 seconds, or after each turn with turn-by-turn live text. That means more requests to a Think server, or more load on a local model during the meeting.
- Turn-by-turn timings are from one scripted call per language on an M4 Max. A slower Mac takes longer.
- Live Radar and Follow-through were built on a hosted model for a hackathon. On a local Think model they are covered by tests with a scripted model, not yet by a real meeting, and speed on a small local model hasn't been measured.